

CONCEPT NOTE · "INTAKETRAINING" WORKSTREAM

Experiential Intake Training — the Coached Intake Simulator

June 2026 · Coastal DME Intake Agent

Train associates to expert level *before* automation — using the same AI that will later automate the routine work — with hands-on simulation instead of desktop SOP reading.

The problem with legacy training

DME/claims intake is still taught the old way: read the SOP deck, watch a few examples, then shadow a senior associate for weeks. It's **passive**, it **doesn't expose people to real variability or edge cases**, there's **no feedback loop**, competency **can't be measured**, and ramp time is long. People learn the rule, not the judgment.

The idea

The production intake agent already knows the right answer for any referral — it extracts fields, resolves ICD conflicts, applies 24 cited CMS rules, and decides route-vs-escalate. **Flip it into a coach.** A trainee does the *real task* on a realistic case; the agent **grades and explains instantly, citing the exact rule.**

The loop: realistic referral packet → trainee extracts ICD / spots conflict / flags gaps / decides route-or-escalate → agent grades each decision and coaches with the CMS citation → difficulty ramps → competency is scored and certified.

What makes it experiential (not a slide deck)

- **Do, don't read** — work real packets, make the real decisions.
- **Curveballs injected** — seeded ICD miscodes, missing auth, laterality, bariatric edge cases — like a flight sim throwing failures.
- **Instant, cited coaching** — every case ends with *why*, tied to the rule (e.g., "post-surgical → S-code not M-code, §I.C.19.a").
- **Judgment reps** — explicitly trains *when to escalate vs. decide* — the exact judgment automation leaves to humans.
- **Gamified & measured** — scoring, levels, a certification gate before live work; manager dashboards on accuracy and error patterns.

Why it beats legacy SOP training

Legacy desktop SOP	Coached simulation
Passive reading	Active doing on real cases
No feedback	Instant, cited coaching every case
Generic rules	Learn from variability + edge cases
Weeks of shadowing	Days to competency
Can't measure readiness	Scored, certified, data-driven

Strategic payoff

- **A bridge to automation, not separate from it** — the same engine that coaches humans now automates the routine cases later. One asset, two uses.
- **De-risks change management** — associates coached *by* the AI trust it and understand it; reframes "AI is replacing me" into "AI made me an expert faster." The #1 adoption blocker, addressed.
- **Raises the human-in-the-loop quality floor** — after automation, humans handle the hard escalations, so you want them expert. Upskilling, not deskilling.
- **Differentiator** — pairing "we'll automate your intake" with "and we'll make your people experts first, using the same AI" is a story competitors don't tell.

Feasibility — it's already prototyped

A working **Coached Intake Simulator** is live in the demo (/panel/training): three curated cases (post-surgical miscode, missing authorization, the bariatric trap); the trainee submits ICD / conflict / gaps / decision, and the agent grades each dimension with cited coaching. It reuses what already exists — referral generators (unlimited cases + ground truth), the agent's KG citations (the coach), the escalation logic (judgment), and the observability layer (competency dashboards).

Next steps: expand the case bank by difficulty tier · add scoring/leaderboard + certification gate · add a manager competency dashboard · optionally a simulated inbound voice case. A great hands-on station for the June 10 workshop.

"Intakettraining" concept note · Coastal DME Intake Agent Platform · June 2026.