

# Design → Demo → Deploy

## ◀ OPERATING PRINCIPLES

- Autonomous where confident, human where not**

The agent never guesses on a clinical or coding decision — it escalates to a specialist.
- Deterministic compliance, not probability**

Coding is validated against 24 cited CMS rules — confirm or reject, never a guess.
- Every decision logged & auditable**

Full trail per episode: reasoning, confidence score, and the exact rule that fired.
- Fits your stack**

API / RPA / file — no rip-and-replace. PHI stays inside your boundary.

## THE DAY — SUMMARY AGENDA

-  **Morning — Demonstrate & Build Trust**

**9:00** Welcome & objectives

**9:30** Live demo — real referrals 

**10:15** Architecture & workflow

**11:00** Trust & human-in-the-loop

**11:40** Integration & knowledge graph 
-  **Afternoon — Prove Fit & Plan Rollout**

**1:10** Security, compliance & enterprise

**1:50** Business case & ROI

**2:45** Deployment & rollout 

**3:25** Roadmap & ecosystem

**3:55** Decisions, owners & next steps

*1 day · 9:00–4:30 · hands-on, not a slideshow*

## WHAT WE’LL PROVE ▶

- The design is sound**

Modular agent + OWL knowledge graph, human-in-the-loop at every flagged exception.
- The demo is real**

A live Claude agent on actual referral PDFs — on your representative data.
- The rollout is feasible**

A concrete 8–12 week path into your environment — owners & dates agreed in the room.
- It’s a platform**

Agent 1 (intake) live today; 7 more agents in the referral-to-payment value chain.

The arc in one line: **Wow → Trust → Fit → Justify → Plan → Commit**